

Challenge: Count Duplicates

Duplicate data occurs when the same row appears more than once in a dataset. These duplicate entries can skew your analysis by overrepresenting certain values, leading to inaccurate statistics, misleading trends, and unreliable results. Detecting and quantifying duplicate rows is a fundamental part of data cleaning, as it helps you understand the extent of the problem and informs your next steps—such as removing or consolidating these duplicates.

```
import pandas as pd

data = {
    "Name": ["Alice", "Bob", "Alice", "Charlie", "Bob", "Alice"],
    "Age": [25, 30, 25, 35, 30, 25],
    "City": ["NY", "LA", "NY", "SF", "LA", "NY"]
}
df = pd.DataFrame(data)
print(df)
```